FOR BIOSTATISTICIANS

Study Operations Knowledge Base — two ways to build it

*A self-updating, queryable knowledge base built automatically from your study Outlook mail — either on **M365 Copilot**, or **Copilot-free on a local Small Language Model**. Same goal, two very different IT and reproducibility profiles.*

The one rule, both ways. This is the *operational / knowledge layer only*. The AI classifies, extracts, routes, and drafts short grounded text — it **never produces a reported number, never reconciles, never decides**. Validated tools (Phoenix WinNonlin, Pinnacle 21, the validated pipeline) and a qualified human own every regulated value; the KB is informational ops support, not a source of record. No PHI / unblinded / participant-level data.

Which path?

An honest read — this is about your tenant and your QA posture, not which model is "better."

CHOOSE THE LOCAL-SLM PATH WHEN...

Copilot-free, on a local frozen SLM

- Reproducibility matters — you can **freeze the model** (tag+quant+SHA-256+seed) into a re-runnable validation artifact.
- You want to **own the reliability controls** (schema-constrained decoding, allowlist validator, human gate) in your own SAS/R code.
- You want to **avoid the Copilot-specific IT frictions** (DLP-connector block, the 01-Nov-2026 AI-Builder credit clock, Copilot Studio licensing + agent-publish gate, cloud egress).
- Honest cost: a small (7–8B) model **hallucinates more** than a frontier model, and it is a **real, qualified instrument** to validate, own, and maintain (its own IT hurdles).

CHOOSE THE M365 COPILOT PATH WHEN...

M365-native, on Copilot + AI Builder

- You are **already entitled**, the DLP question is sorted (the three connectors in one group), and the content is genuinely de-identified ops-only.
- You want it **M365-native and fast** — for the ops/drafting layer Copilot is genuinely capable, at a few dollars/month.
- You **cannot freeze the model** (no reproducible validation artifact), and you carry the DLP / credit-clock / licensing / agent-publish frictions — configuration states you can resolve, but real.

And the option that ships Monday: Outlook + SharePoint + native search + a light template — *no model at all* — captures a surprising share of the value with zero AI entitlement and zero validation burden. If that is enough, it is a fair reason to *defer* standing up either AI path; the model's marginal value must clear its own cost.

Track 1 · The Copilot-free local-SLM path (new)

Local SLM operating wiki

The full Copilot-free build: a local scheduled Graph job, an on-prem frozen SLM doing schema-constrained extraction (LOCALMIND), a human-gated parse-guard, SharePoint store, a local RAG "ask the KB", and a deterministic-first weekly digest — with the honest reliability story and the IT-readiness contrast kept in the headline.

Open wiki (HTML) PDF

The engine & the IT playbook (reused)

The LOCALMIND `%slm_*` library (loopback-guard, schema-constrained, allowlist-validator) that does the extract+validate+gate, the on-device worked example, and the IT/Security/QA enablement runbook with the hurdles + open-weight-licence landscape to hand to IT.

On-device SLM wiki IT runbook ▶ Example

Track 2 · The M365 Copilot path

Biostatistics × Copilot operating wiki

The original: the self-updating Outlook→Power Automate→AI Builder→SharePoint→Copilot-agent pipeline, with build screencasts (the flow grows action-by-action), the worked example, the prompt library, and governance.

Copilot IT readiness

What to ask IT for the Copilot path (the DLP-connector group, the AI-Builder credit clock, the agent-publish control), and the ships-Monday no-AI fallback.

Why the AI step is kept reproducible — the evidence

Field evidence — Copilot on a real QC task

A verbatim, page-cited record of the M365 Copilot Analyst Agent on an SDTM→ADaM→TLF QC task: empty "FINAL" files, walked-back "submission-ready" claims, an unstable error rate, a fabricated appendix — and its own concession "you shouldn't [use this]." *Read fairly: one high-stakes QC session, not the ops-KB task — it shows why to keep anything you must reproduce off a non-pinnable cloud model.*

Open exhibit (HTML) PDF

Copilot trial-ops feasibility

The 75-automation re-assessment: where Copilot is genuinely feasible (19/75 fully, most buildable) and where it falls short — the gaps tracing to no-on-prem-model + no-model-freeze, exactly what the local path closes.

Open assessment PDF

Both paths are the operational layer only; validated tools & a qualified human own every reported number. A frozen on-prem model is reproducible, not magically reliable — a small model hallucinates, and the local path is a real validation/ownership commitment. The wikis are honestly-labelled designs; capability is established only by your own OQ validation.